

Dockercon EU 2017

Networking Workshop

Mark Church, Technical Account Manager @ Docker

Nico Kabar, Solutions Architect @ Docker



Agenda

1. Fundamentals & Network Drivers
2. Bridge Driver
3. Overlay Driver
4. Network Services (DNS, Internal and External Load Balancing, Publishing)
5. Lab I
6. BREAK (15 minutes)
7. MACVLAN Driver
8. Network Design & Best Practices
9. Network Troubleshooting
10. Lab II

Docker Networking Fundamentals

Docker Networking Design Philosophy

Put Users First

Developers and
Operations

Plugin API Design

Batteries included
but removable

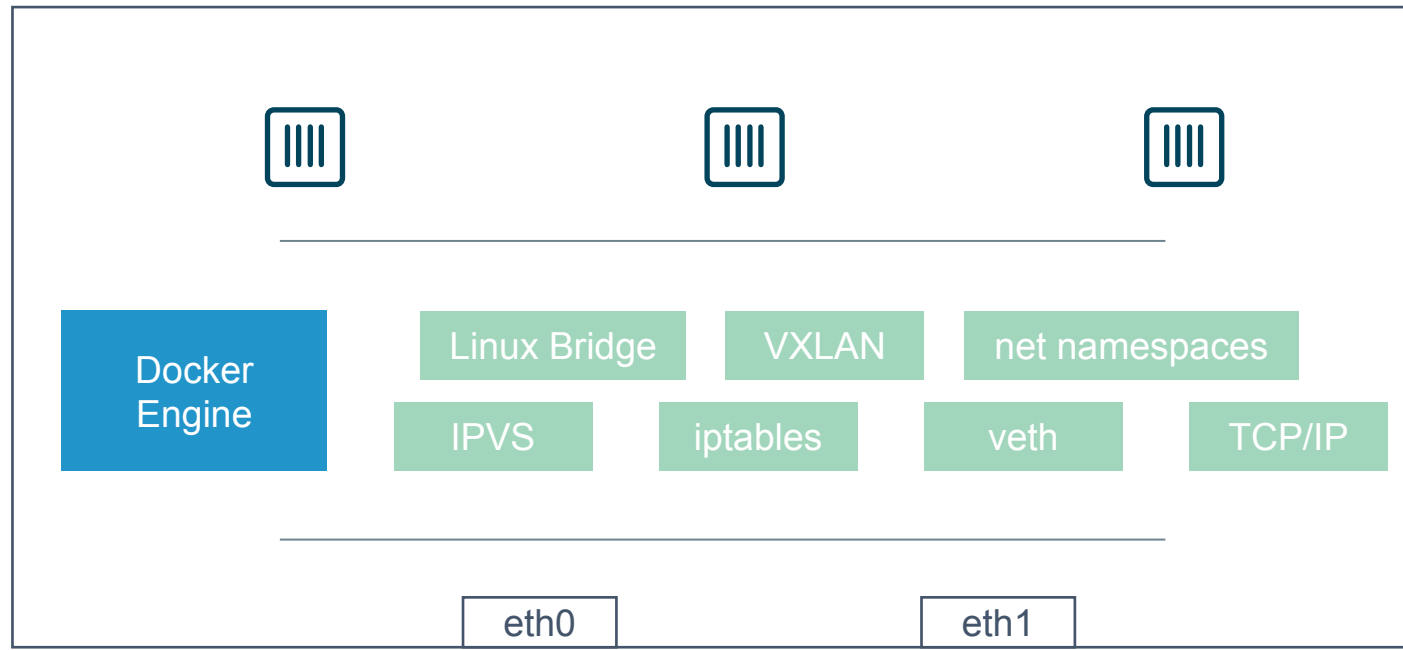
Docker Networking *is* Linux (and Windows) Networking

Host

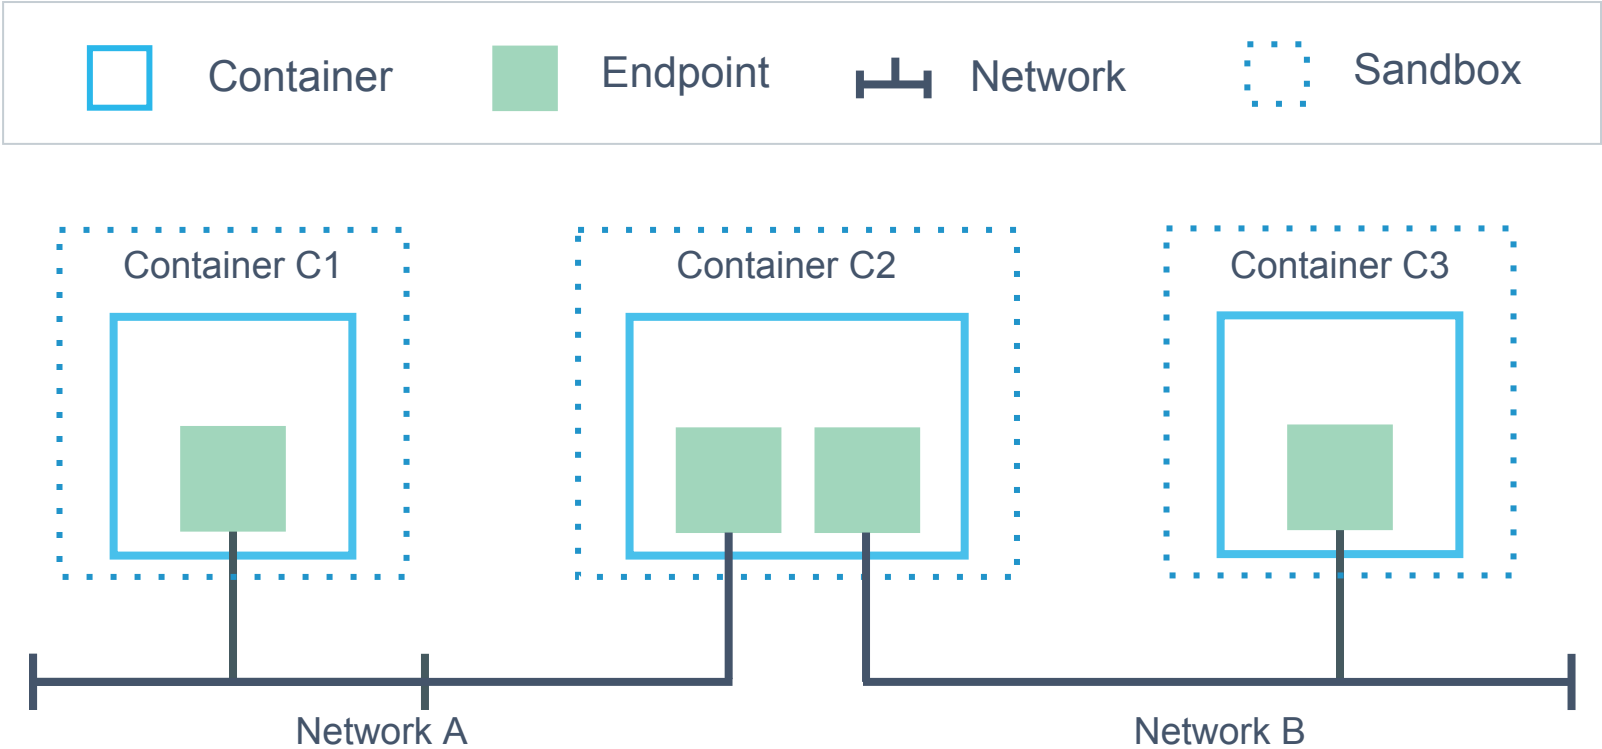
User Space

Kernel

Devices



Containers and the CNM



Native Docker Networking Drivers

```
$ docker info
```

```
...
```

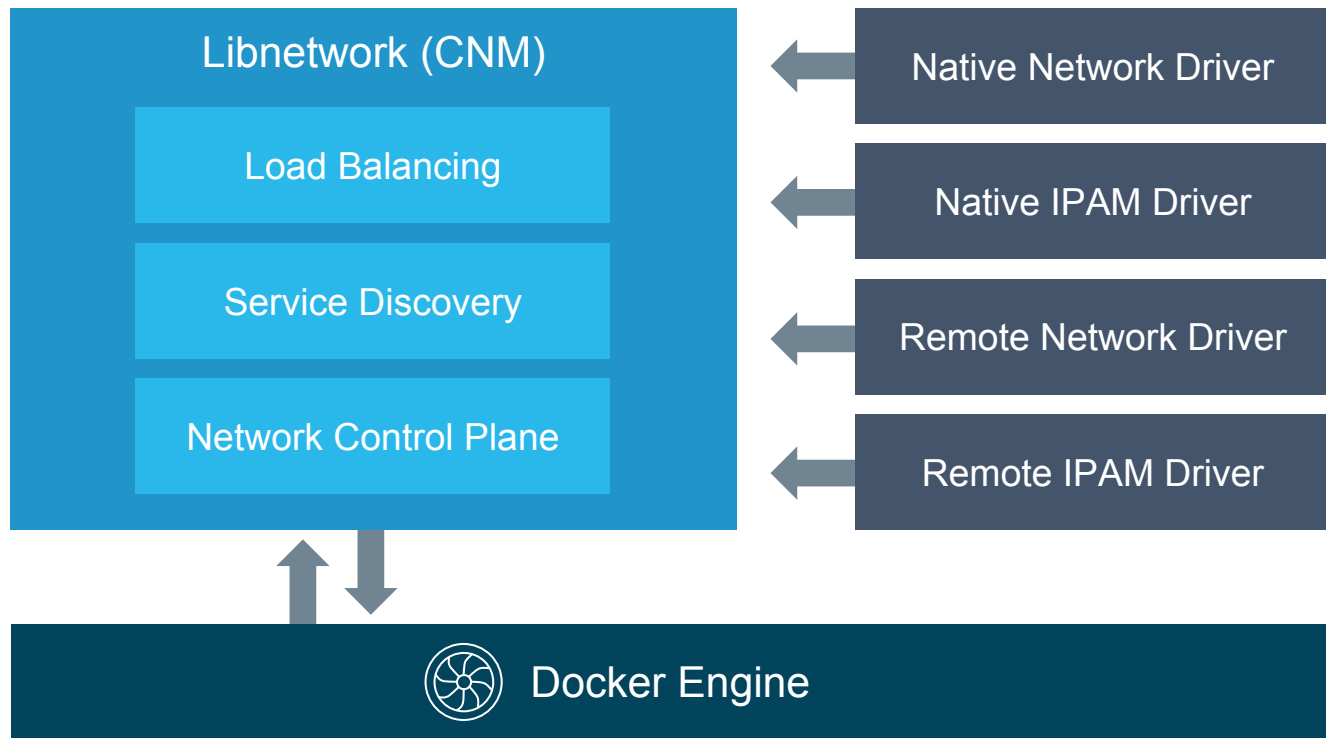
```
Plugins:
```

```
Volume: local
```

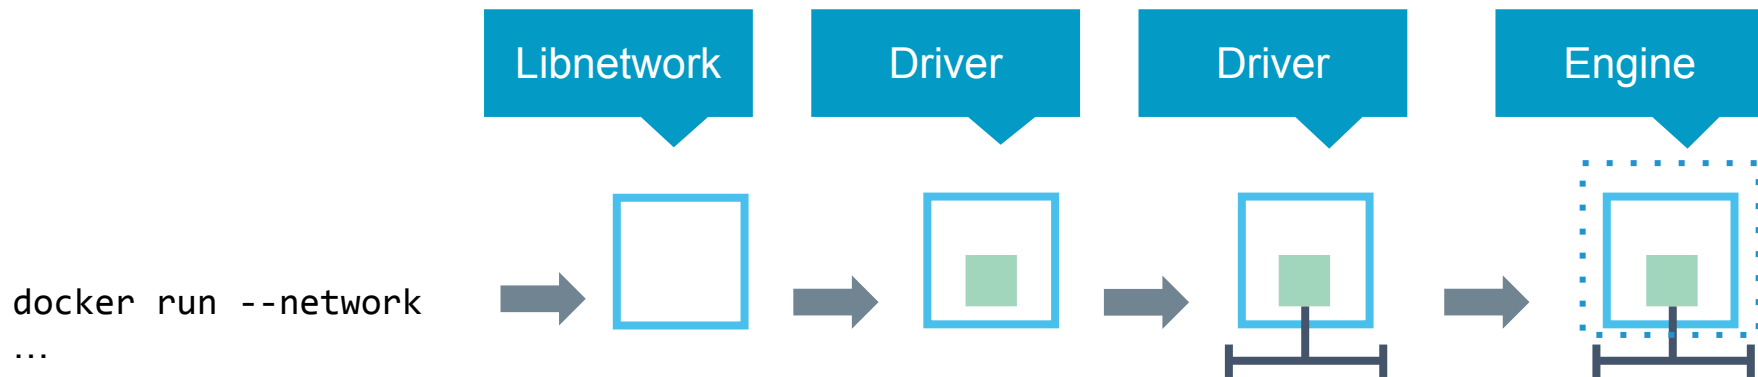
```
Network: bridge host ipvlan macvlan null overlay
```

```
...
```

Libnetwork Architecture



Networks and Containers




Docker Networking on Linux

- The Linux kernel has extensive networking capabilities (TCP/IP stack, VXLAN, DNS...)
- Docker networking utilizes many Linux kernel networking features (network namespaces, bridges, iptables, veth pairs...)
- Linux bridges: L2 virtual switches implemented in the kernel
- Network namespaces: Used for isolating container network stacks
- veth pairs: Connect containers to container networks
- iptables: Used for port mapping, load balancing, network isolation...


Docker Windows Networking

Container Networking Model



Linux Networking with Containers


- Namespaces are used extensively for container isolation
- Host network namespace is the default namespace
- Additional network namespaces are created to isolate containers from each other



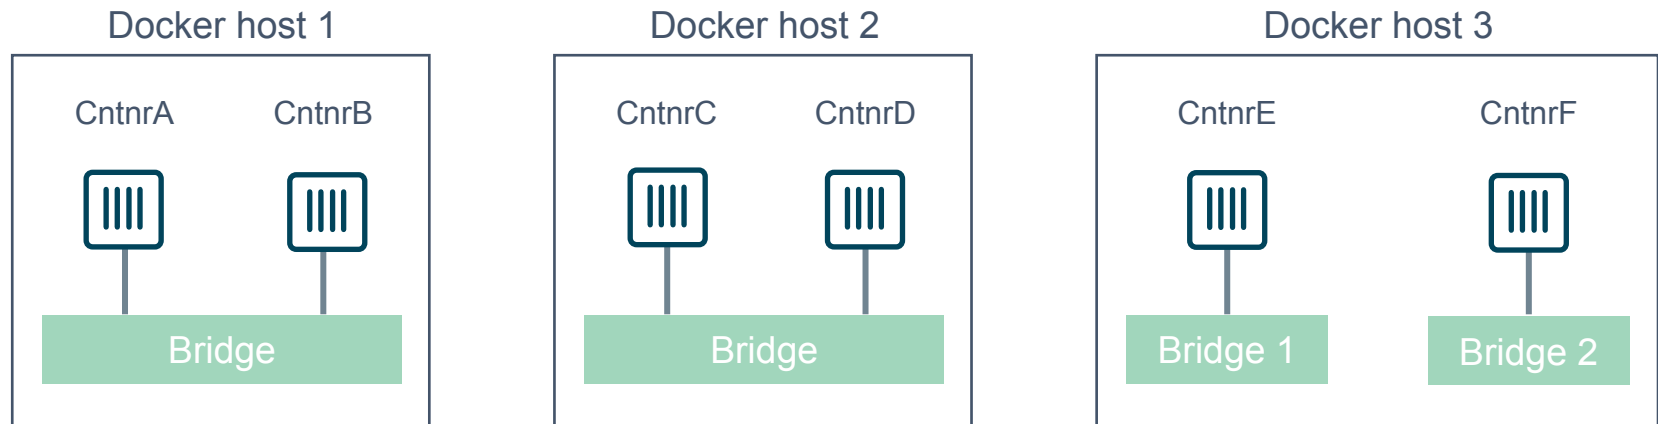
Bridge Driver

Bridge Driver in Detail

- The bridge created by the bridge driver for the pre-built bridge network is called docker0
- Each container is connected to a bridge network via a veth pair which connects between network namespaces
- Provides single-host networking
- External access requires port mapping




What is Docker Bridge Networking?



Containers on different **bridge** networks cannot communicate

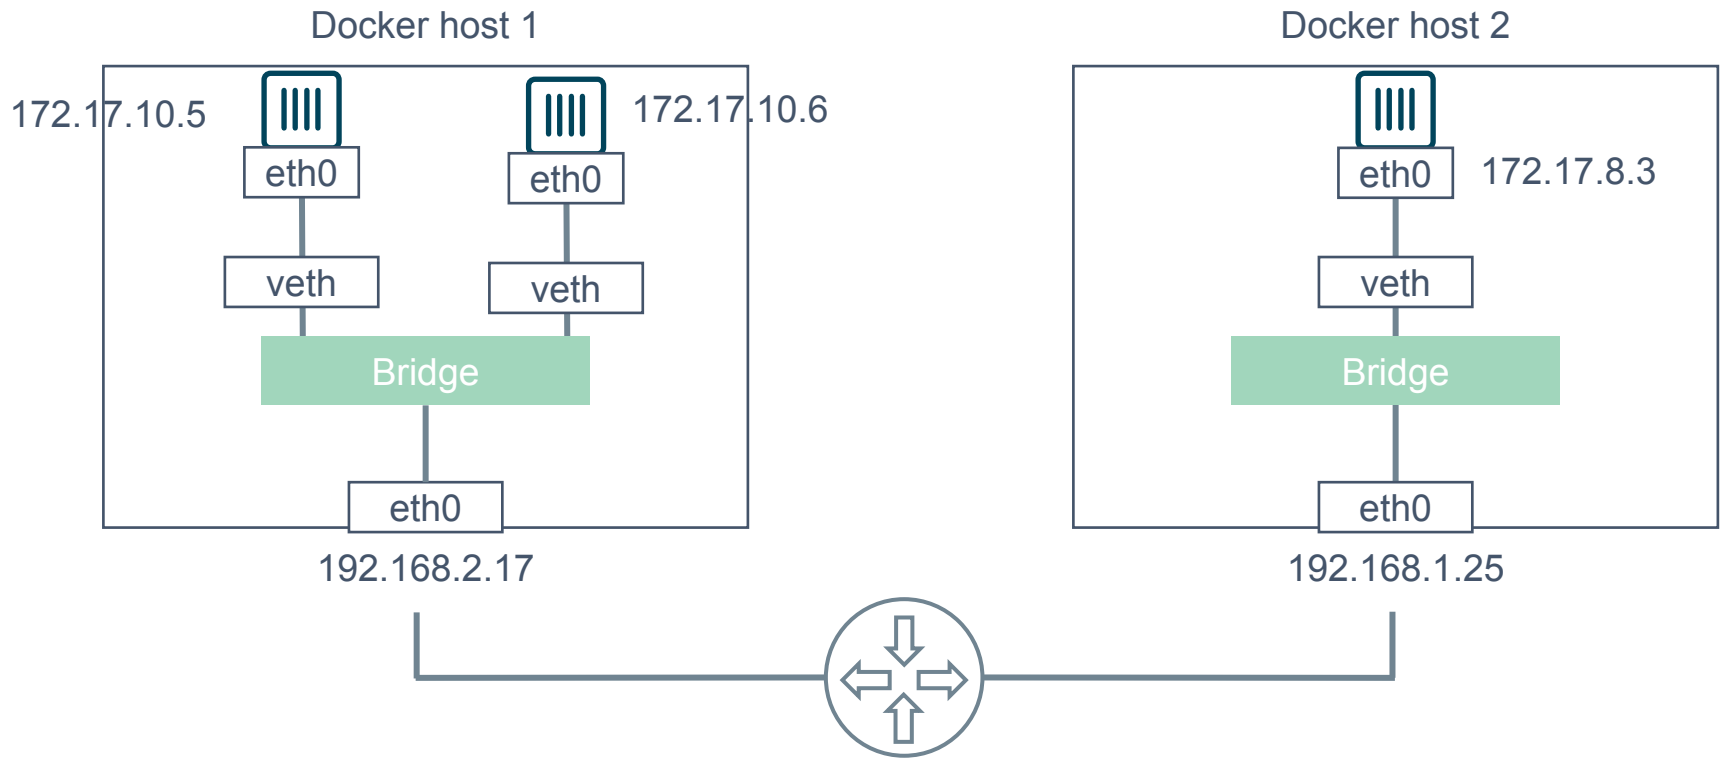
Docker Bridge Networking and Port Mapping



Host port Container port

```
$ docker run -p 8080:80 ...
```



Bridge Mode Data Flow



Overlay Driver


What is Docker Overlay Networking?

The **overlay** driver enables simple and secure **multi-host** networking




All containers on the **overlay** network can communicate!

Building an Overlay Network (High level)




Docker Overlay Networks and VXLAN

- The **overlay** driver uses VXLAN technology to build the network
- A **VXLAN tunnel** is created through the **underlay network(s)**
- At each end of the tunnel is a VXLAN tunnel end point (**VTEP**)
- The **VTEP** performs encapsulation and de-encapsulation
- The **VTEP** exists in the Docker Host's network namespace




Overlay Network Encryption with IPsec



Docker Network Control Plane

Docker networking

- Provides portable application services
 - Service-Discovery
 - Load-Balancing
- Built-in and pluggable network drivers
 - Overlay, macvlan, bridge
 - Remote Drivers / Plugins
- Built-in Management plane
 - API, CLI
 - Docker Stack / Compose
- Built-in distributed control plane
 - Gossip based
- Encrypted Control & Data plane




Gossip

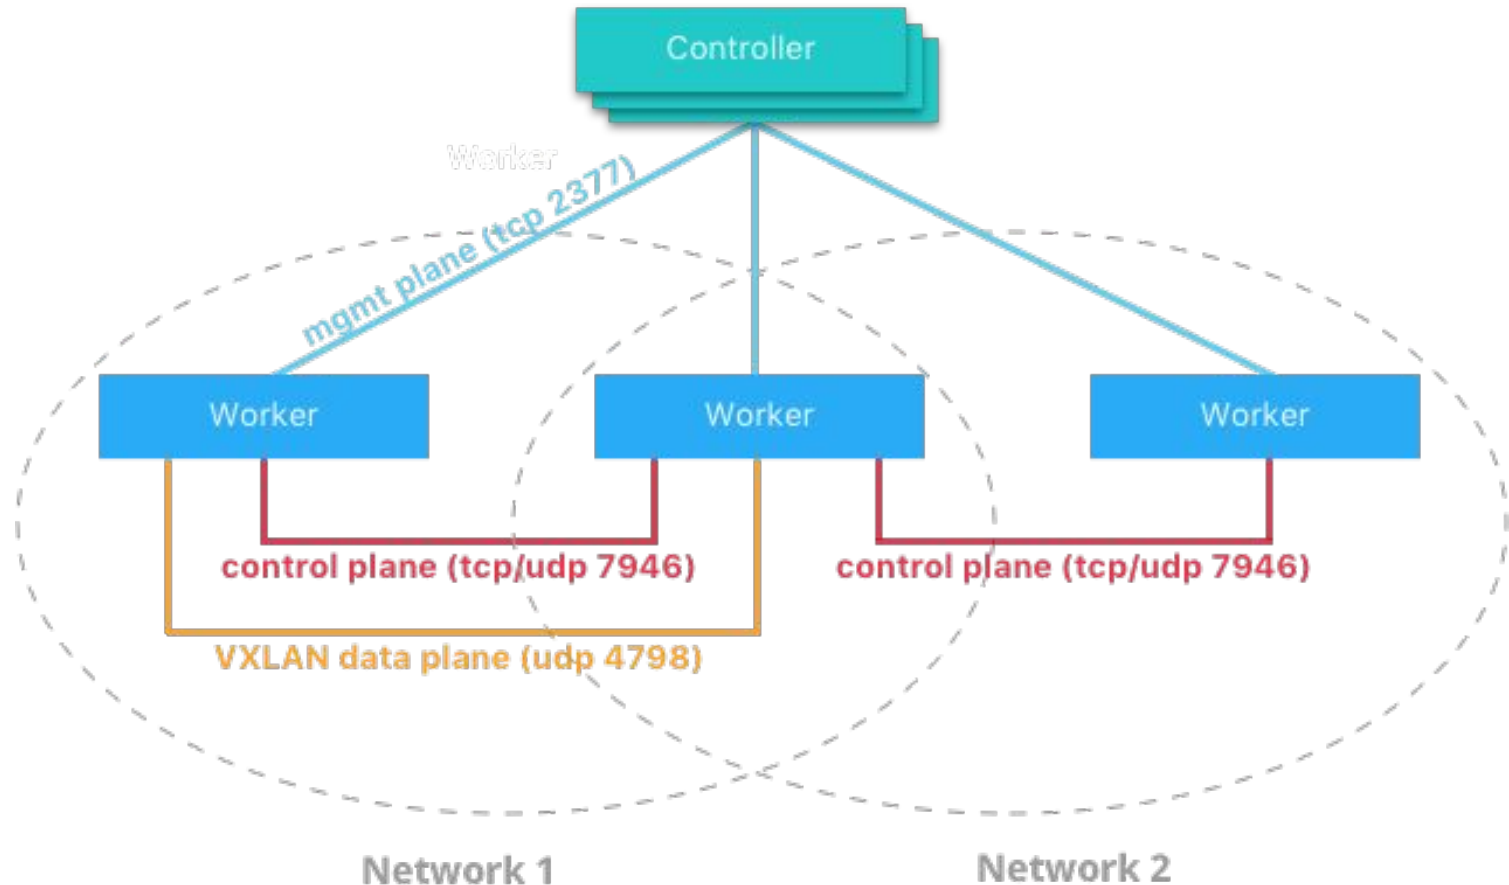
- **Eventually consistent**
- **State dissemination through de-centralized events**
 - Service Registration
 - Load-Balancer configs
 - Routing states
- **Fast convergence**
 - $\sim O(\log n)$
- **Highly scalable**
- **Continues to function even if all managers are Down**

Decentralized Event Propagation

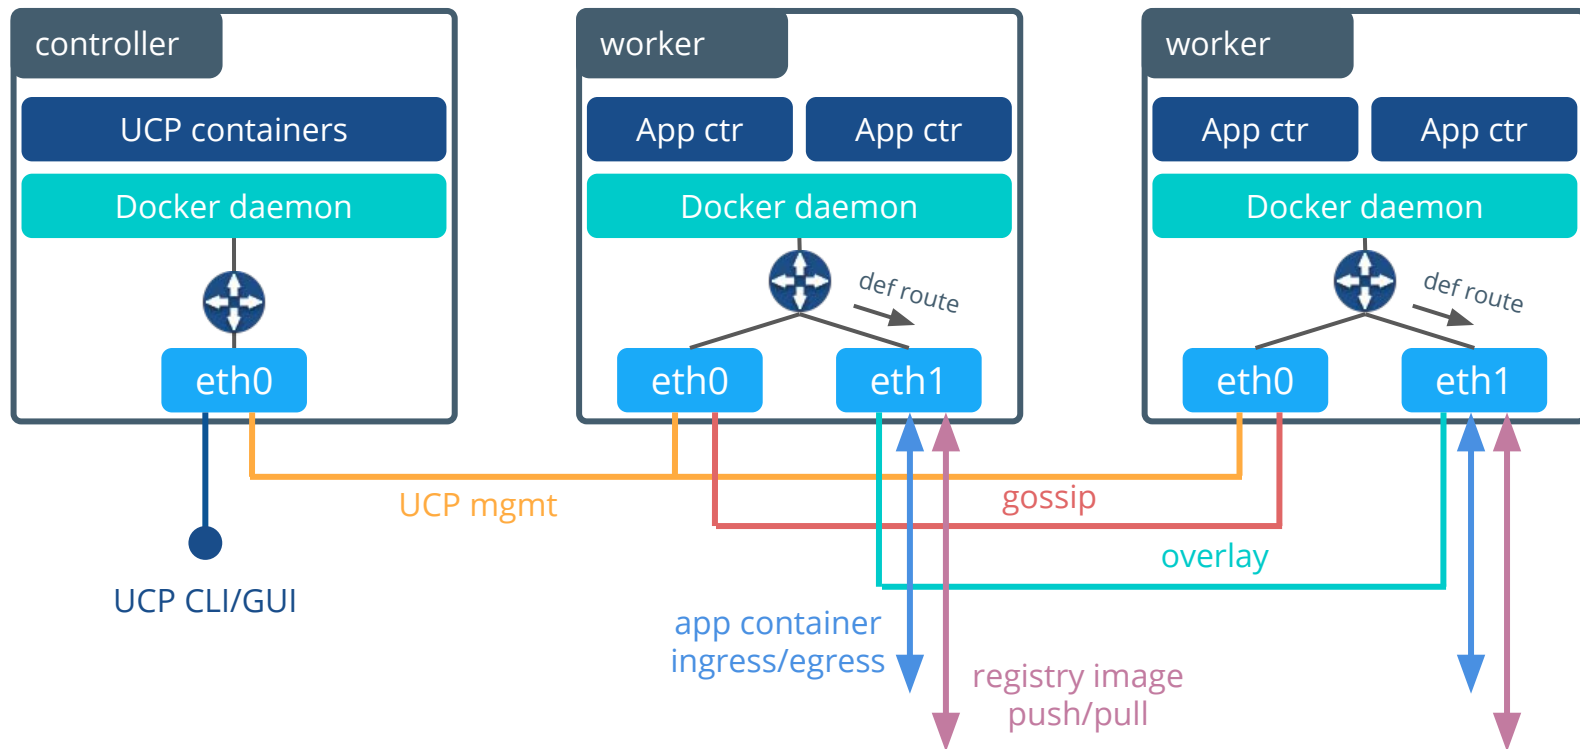
Swarm Scope Gossip



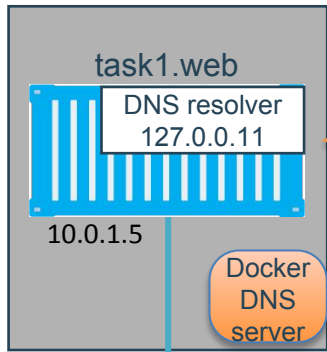
Mgmt, Control, and Data Plane in Practice



Traffic Flows in Detail




Worker1




Gossip

Worker2



Gossip

Worker3



demo_frontend overlay network (vxlan-id 4097)

demo_backend overlay network (vxlan-id 4098)

Service Discovery states

```
web 10.0.1.4 (vip)
app 10.0.1.8 (vip)
task1.web 10.0.1.5
task2.web 10.0.1.6
task1.app 10.0.1.9
```

Routing states

```
10.0.1.6 : {Worker2, 4097}
10.0.1.9 : {Worker2, 4097}
```

Service Discovery states

```
web 10.0.1.4 (vip)
app 10.0.1.8 (vip)
task1.web 10.0.1.5
task2.web 10.0.1.6
task1.app 10.0.1.9
```

Routing states

```
10.0.1.5 : {Worker1, 4097}
```

Service Discovery states

```
db 10.0.2.4 (vip)
app 10.0.2.8 (vip)
task1.db 10.0.2.5
task1.app 10.0.2.6
```

Routing states

```
10.0.2.5 : {Worker3, 4098}
```

Service Discovery states

```
db 10.0.2.4 (vip)
app 10.0.2.8 (vip)
task1.db 10.0.2.5
task1.app 10.0.2.6
```

Routing states


```
10.0.2.6 : {Worker2, 4098}
```

Docker Network Services

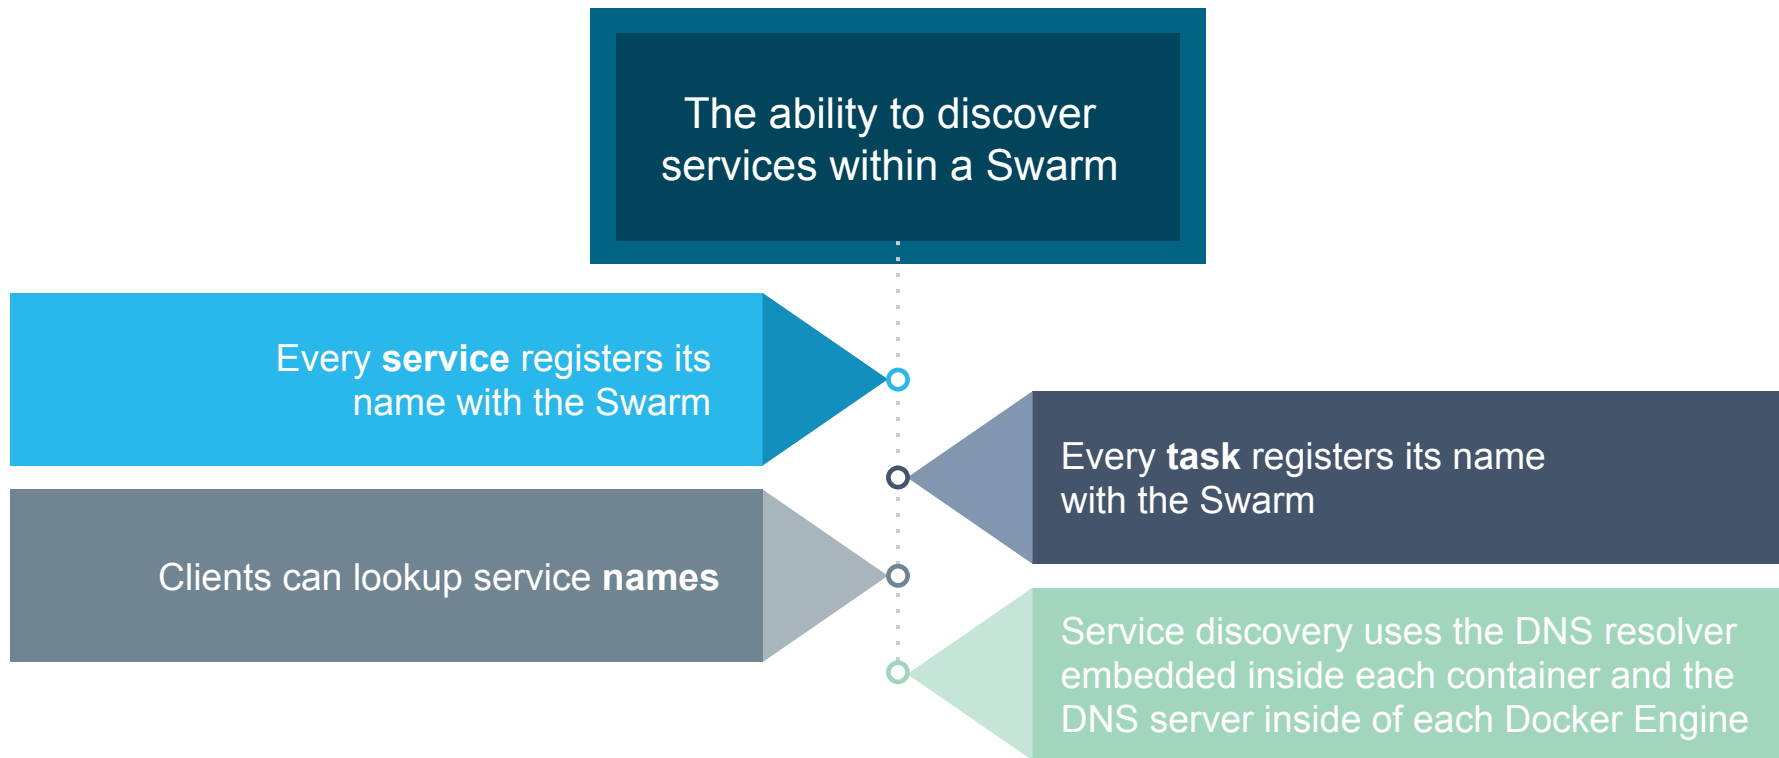
SERVICE REGISTRATION, SERVICE DISCOVERY, AND LOAD BALANCING

Docker Networking

Service Discovery



What is Service Discovery?



Service Discovery Details

Service and task registration is automatic and dynamic



Name-IP-mappings stored in the Swarm KV store

Resolution is network-scoped

Container DNS and Docker Engine DNS used to resolve names

- Every container runs a local DNS resolver (127.0.0.1:53)
- Every Docker Engine runs a DNS service

Docker Stack Deploy

```
$ docker stack deploy -c d.yml demo
```

```
Creating network demo_frontend
```

```
Creating network demo_backend
```

```
Creating service demo_web
```

```
Creating service demo_app
```

```
Creating service demo_db
```

- **Swarm scope** - network resources that are owned and managed centrally by the controllers
- **Local scope** - network resources that are owned and managed by the worker node

Local scope - bridge, macvlan, host

Swarm scope - overlay

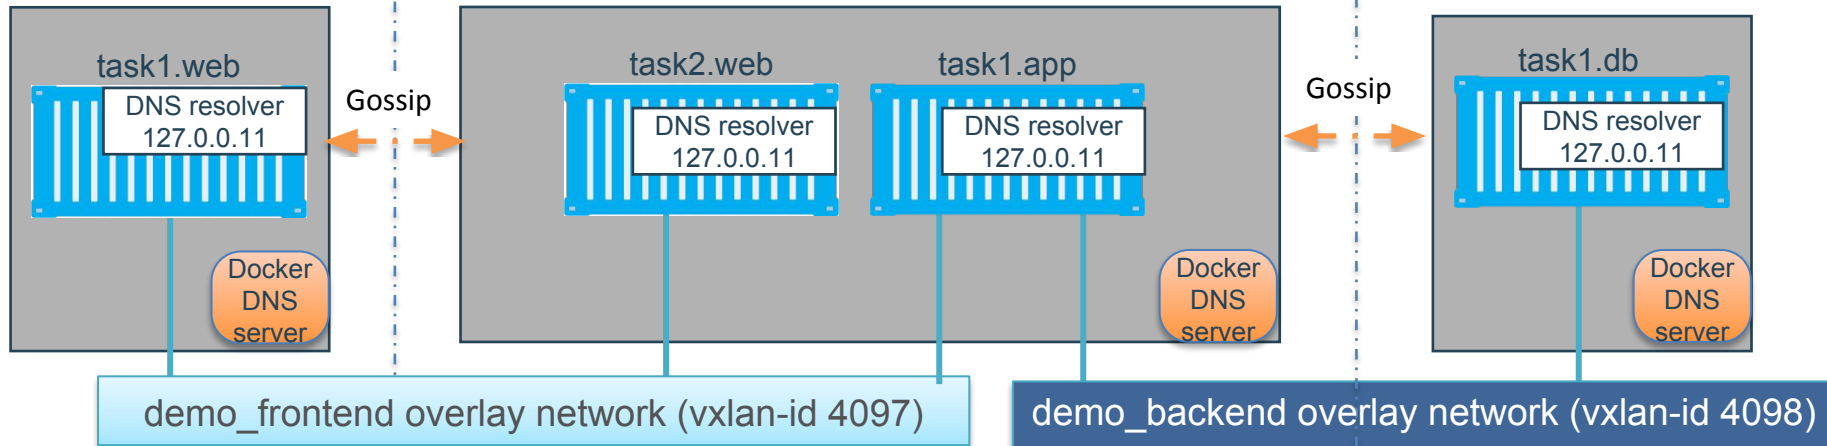
- Manager only operation
- Reserves network resources at mgmt plane such as subnet and vxlan-id. No impact to the data-plane yet.

- Tasks Scheduled to swarm workers
- Network scoped Service Registration on Docker DNS server
 - Service name -> VIP
 - Task name -> Task IP
 - tasks.Service-Name -> All Task IPs
- Exchange SD & LB states via Gossip
- **Prepare Data-plane***
- Call Driver APIs and exchange driver states via Gossip

Worker1

Worker2

Worker3



Service Discovery states

```
web 10.0.1.4 (vip)
app 10.0.1.8 (vip)
task1.web 10.0.1.5
task2.web 10.0.1.6
task1.app 10.0.1.9
```

Routing states

```
10.0.1.6 : {Worker2, 4097}
10.0.1.9 : {Worker2, 4097}
```

Service Discovery states

```
web 10.0.1.4 (vip)
app 10.0.1.8 (vip)
task1.web 10.0.1.5
task2.web 10.0.1.6
task1.app 10.0.1.9
```

Routing states

```
10.0.1.5 : {Worker1, 4097}
```

Service Discovery states

```
db 10.0.2.4 (vip)
app 10.0.2.8 (vip)
task1.db 10.0.2.5
task1.app 10.0.2.6
```

Routing states

```
10.0.2.5 : {Worker3, 4098}
```

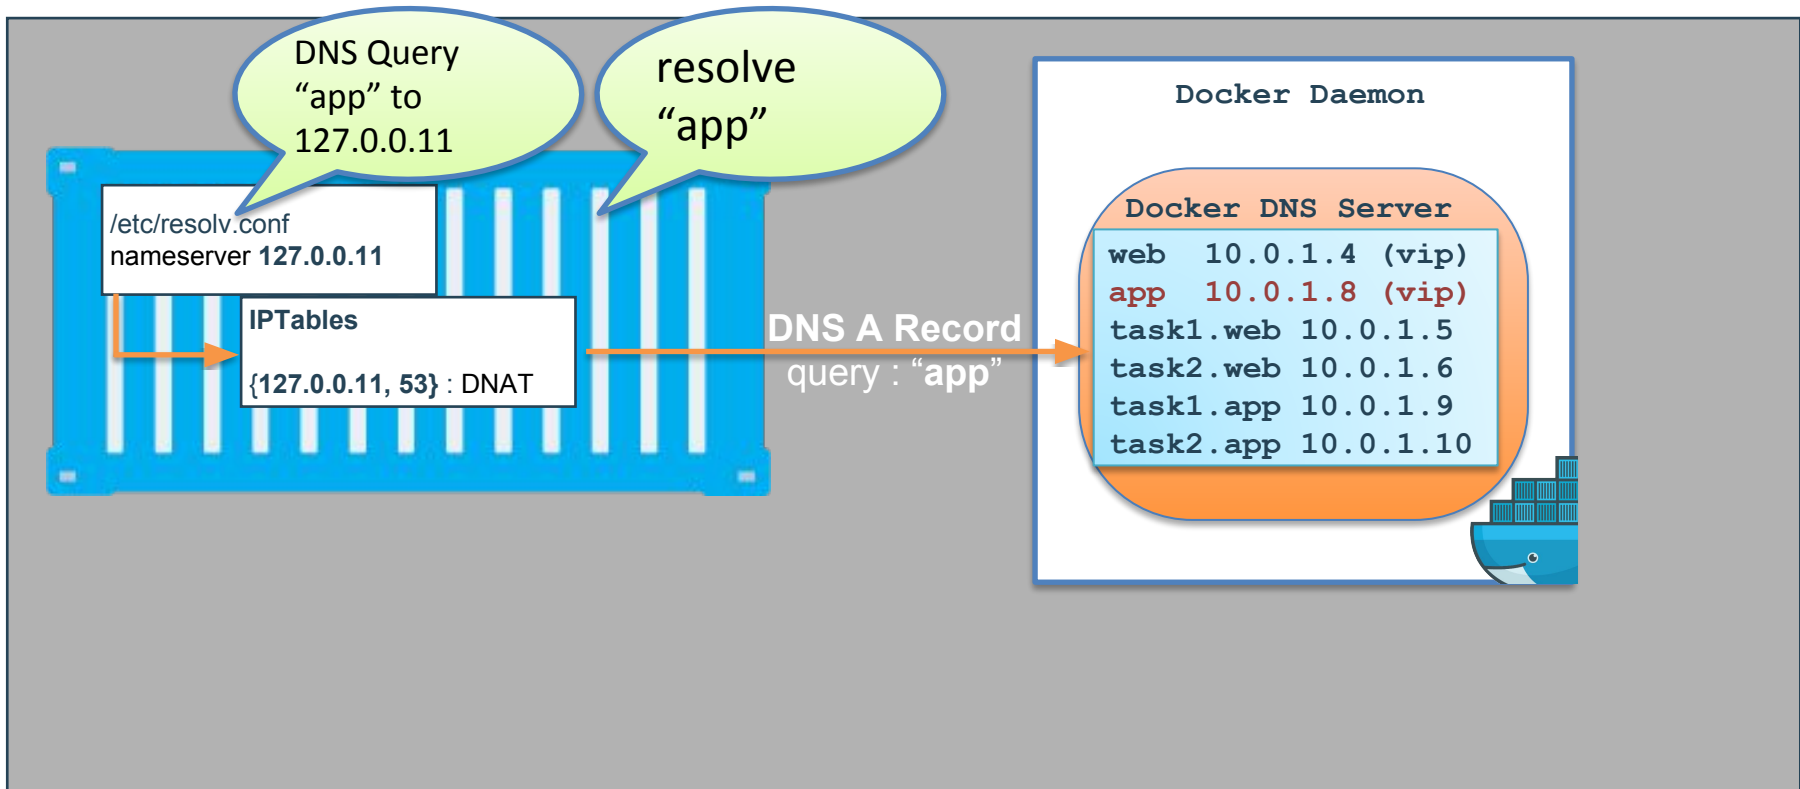
Service Discovery states

```
db 10.0.2.4 (vip)
app 10.0.2.8 (vip)
task1.db 10.0.2.5
task1.app 10.0.2.6
```

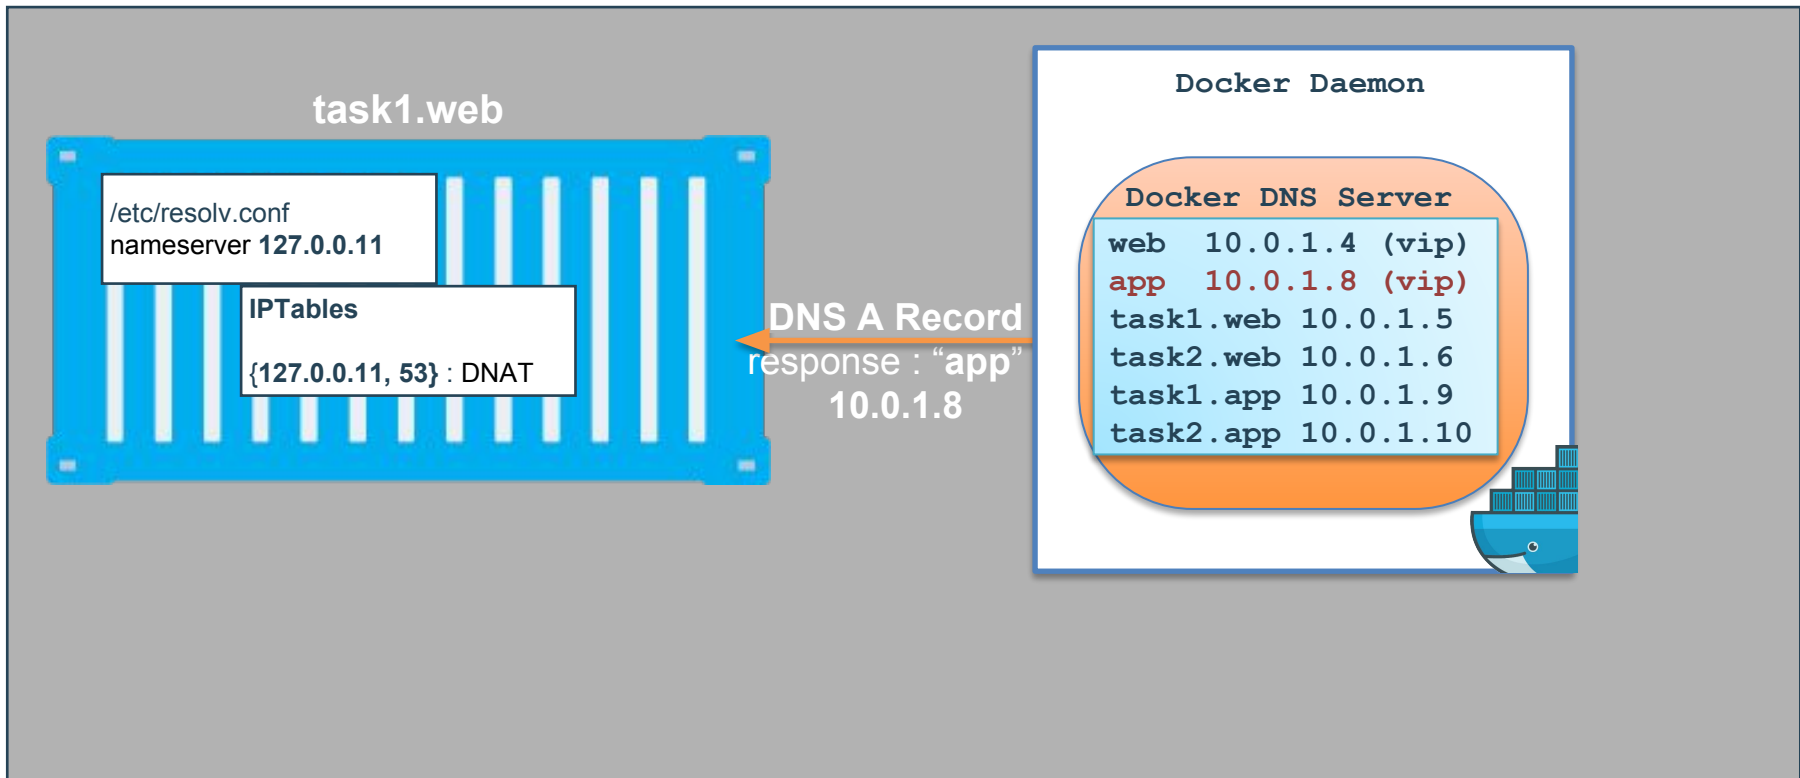
Routing states

```
10.0.2.6 : {Worker2, 4098}
```

Dissecting DNS Lookups




Dissecting DNS Lookups



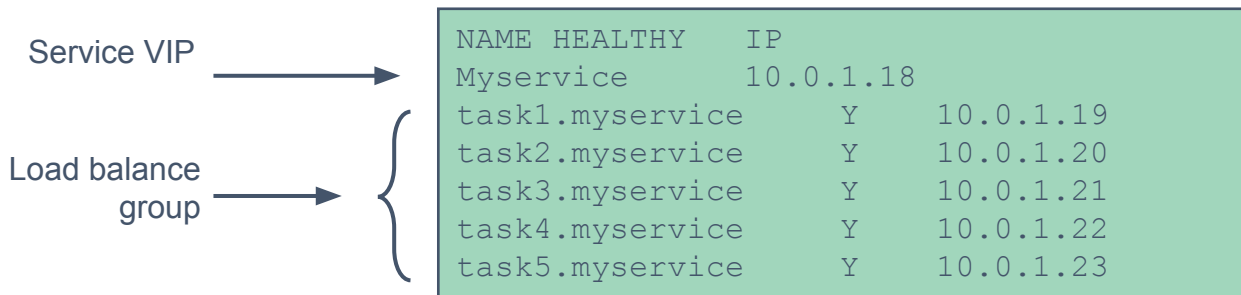
Docker Networking

Load Balancing



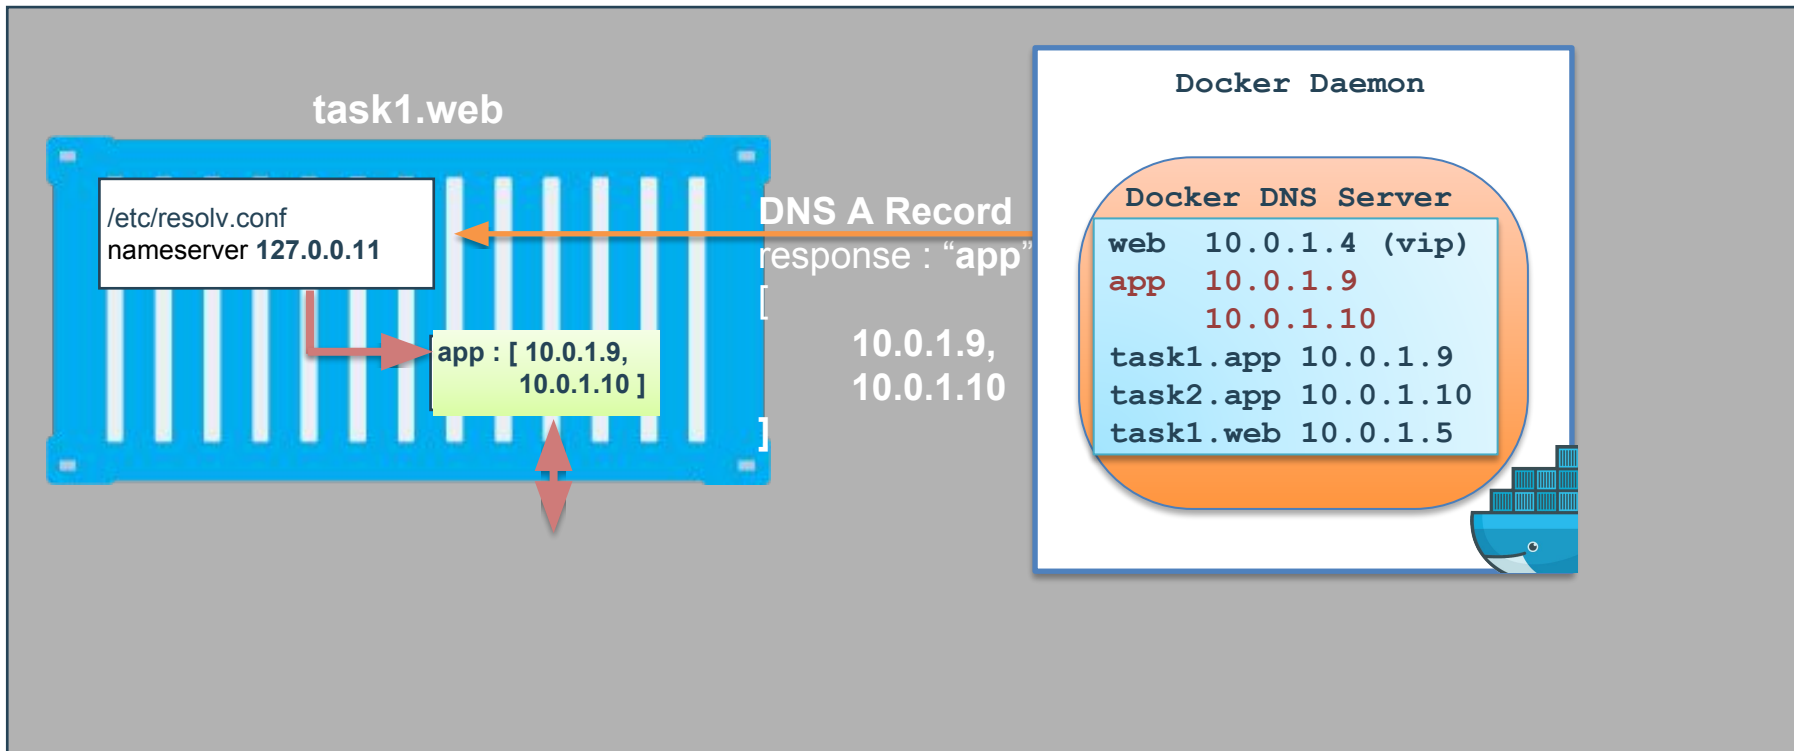
Internal LB: Service Virtual IP (VIP) Load Balancing

- Every **service** gets a **VIP** when it's created (stays with the service for its entire life)
- Lookups against the VIP get load-balanced across all **healthy tasks** in the service
- Behind the scenes it uses Linux kernel **IPVS** to perform transport layer load balancing
- `docker inspect <service>` (shows the service VIP)



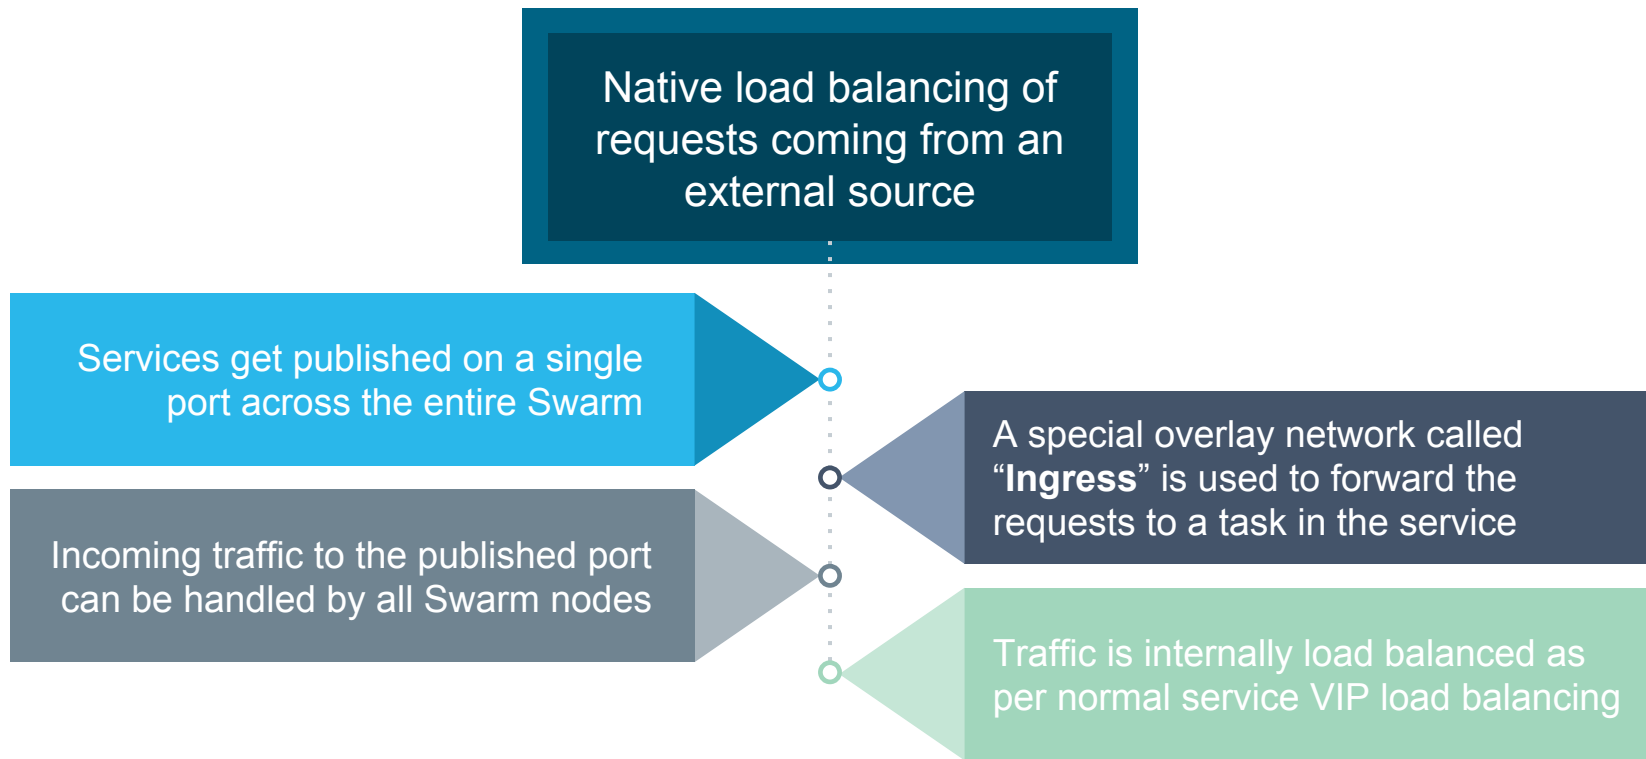
Internal LB: DNS RR Load Balancing

`docker service create --name=app --endpoint-mode=dns-rr demo/my-app`




Publishing Services

What is the Routing Mesh?

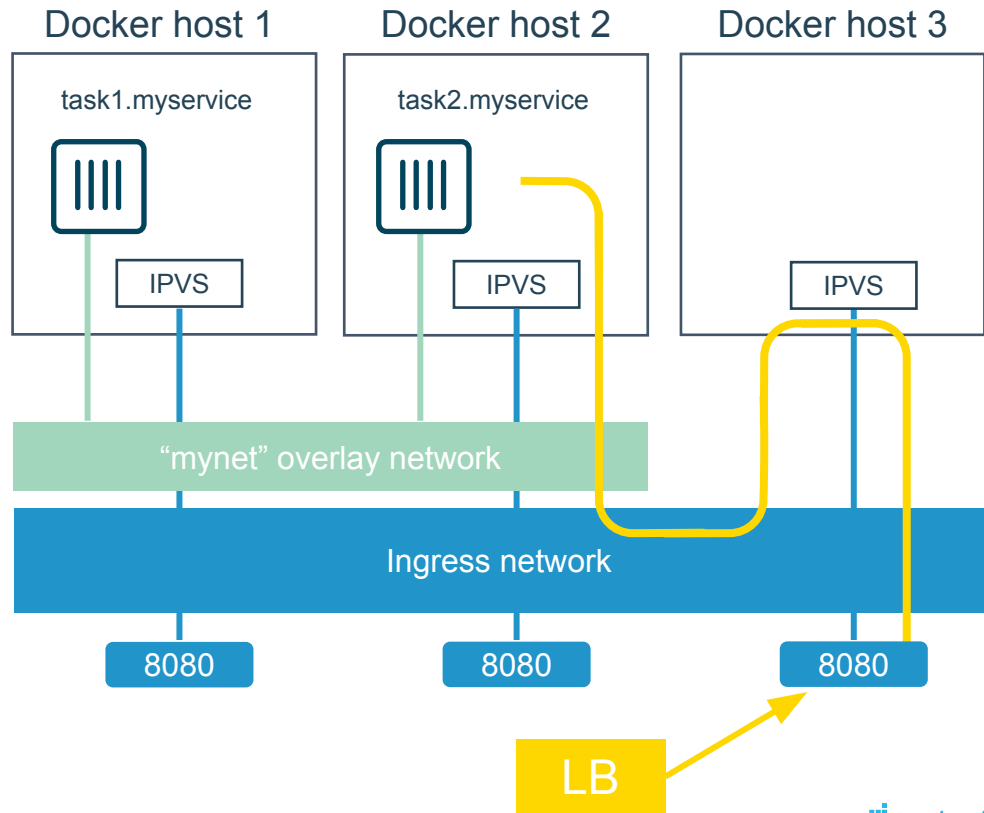


Publish Mode: Host vs Ingress

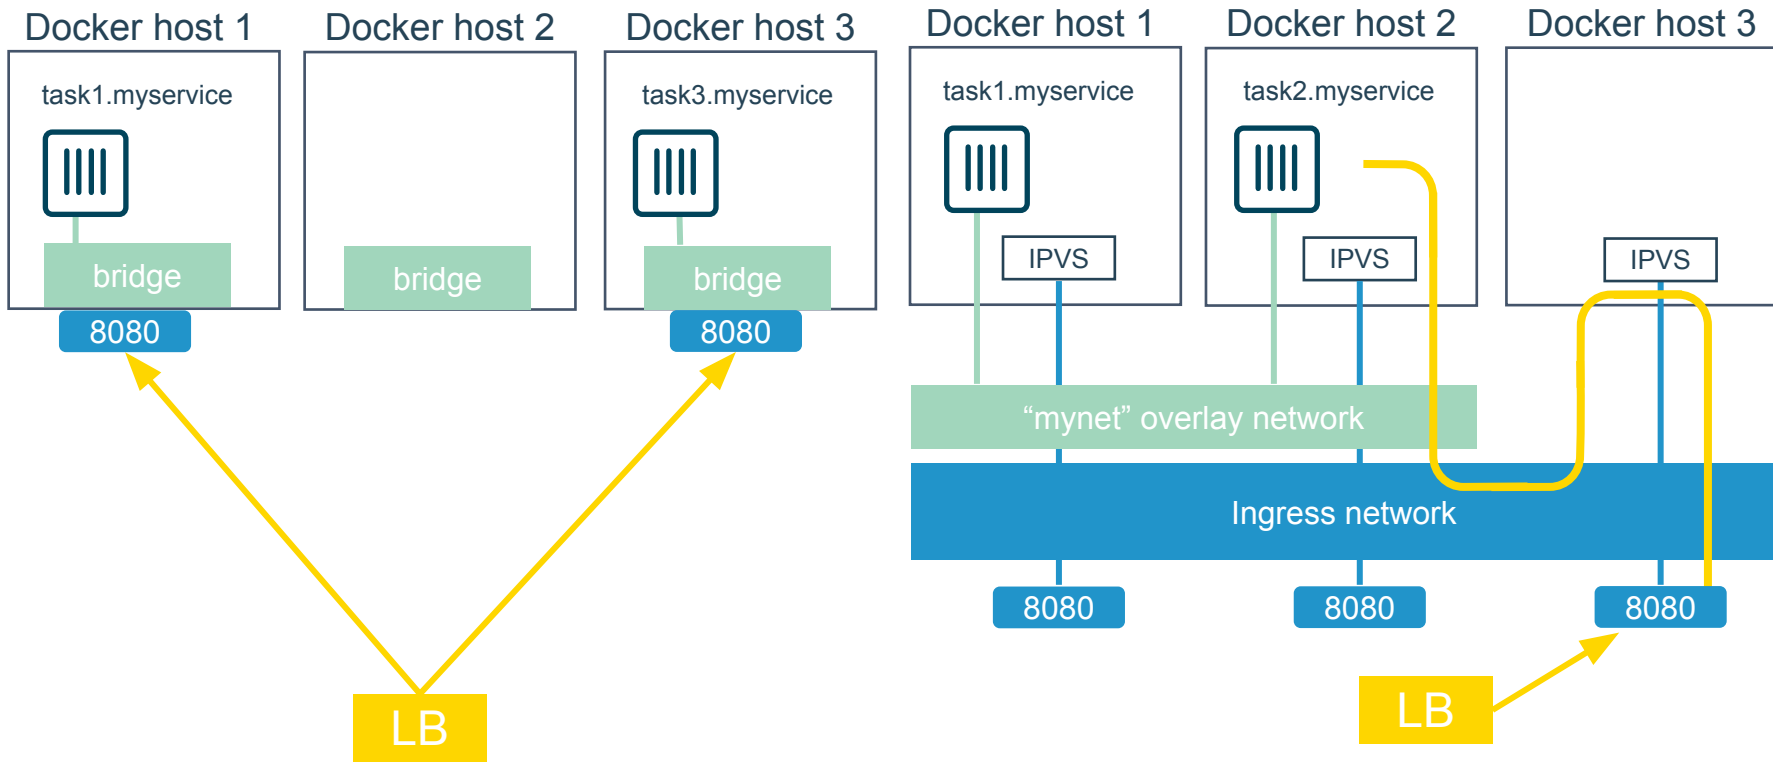


External LB: Ingress Routing Mesh

1. Three Docker hosts
2. New service with 2 tasks
3. Connected to the **mynet** overlay network
4. Service published on port 8080 swarm-wide
5. External LB sends request to Docker host 3 on port 8080
6. Routing mesh forwards the request to a healthy task using the ingress network



Publish Mode: Host vs Ingress




Load Balancing External Requests

ROUTING MESH


Routing Mesh Example

1. Three Docker hosts
2. New service with 2 tasks
3. Connected to the **mynet** overlay network
4. Service published on port 8080 swarm-wide
5. External LB sends request to Docker host 3 on port 8080
6. Routing mesh forwards the request to a healthy task using the ingress network




Routing Mesh Example

1. Three Docker hosts
2. New service with 2 tasks
3. Connected to the **mynet** overlay network
4. Service published on port 8080 swarm-wide
5. External LB sends request to Docker host 3 on port 8080
6. Routing mesh forwards the request to a healthy task using the ingress network



Host Mode vs Routing Mesh



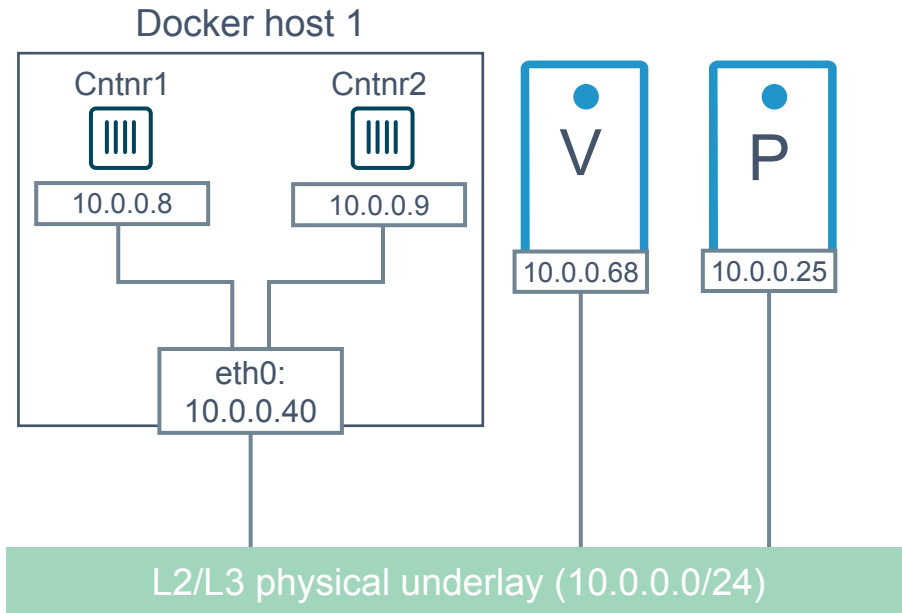
Lab I

BREAK

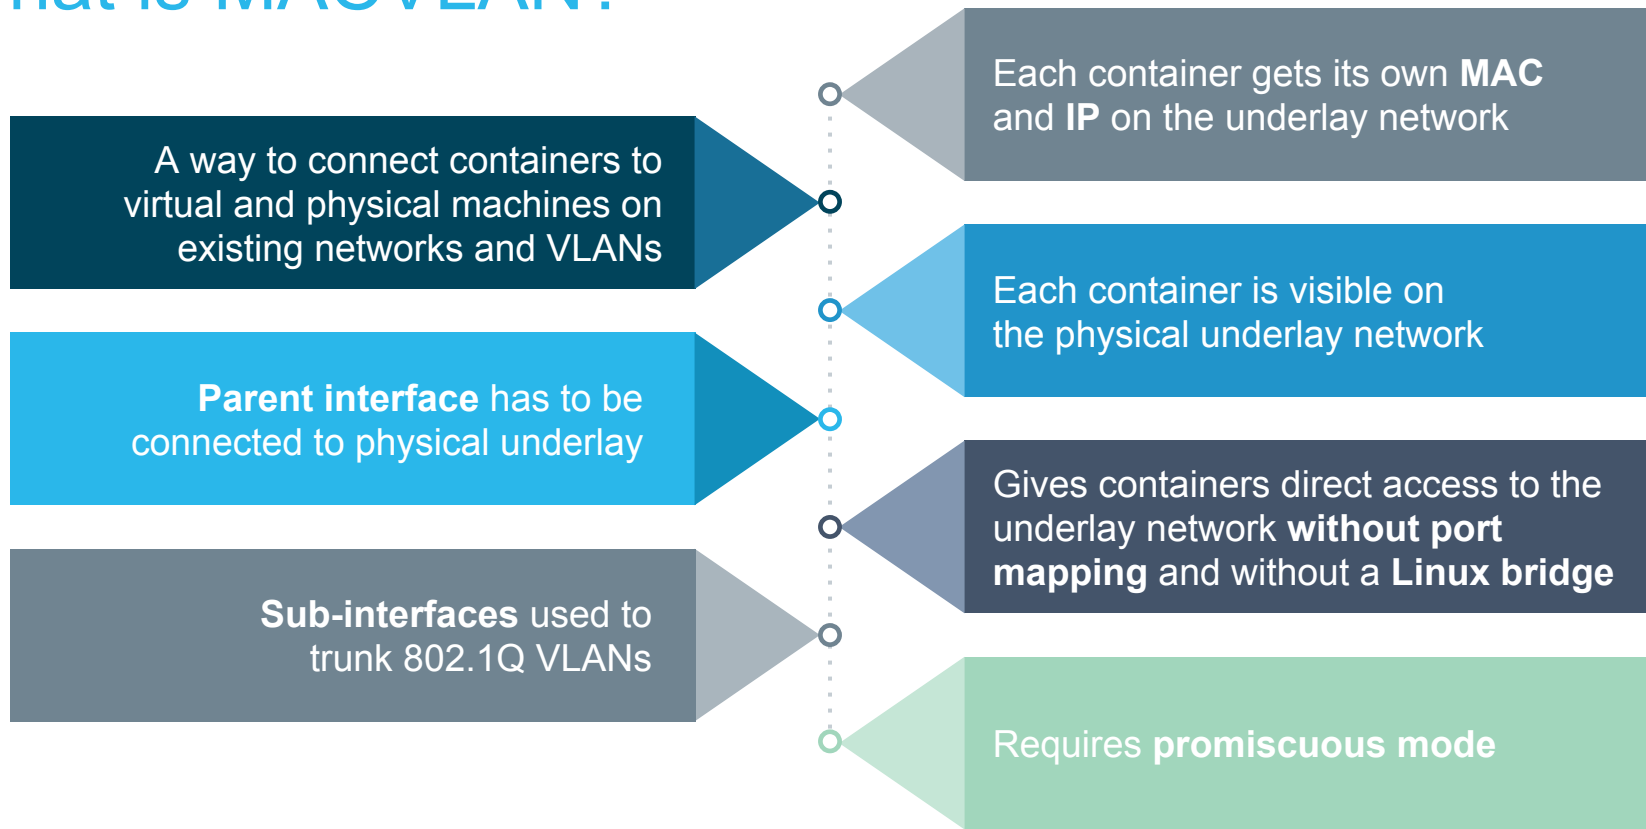
MACVLAN Driver

What is MACVLAN?

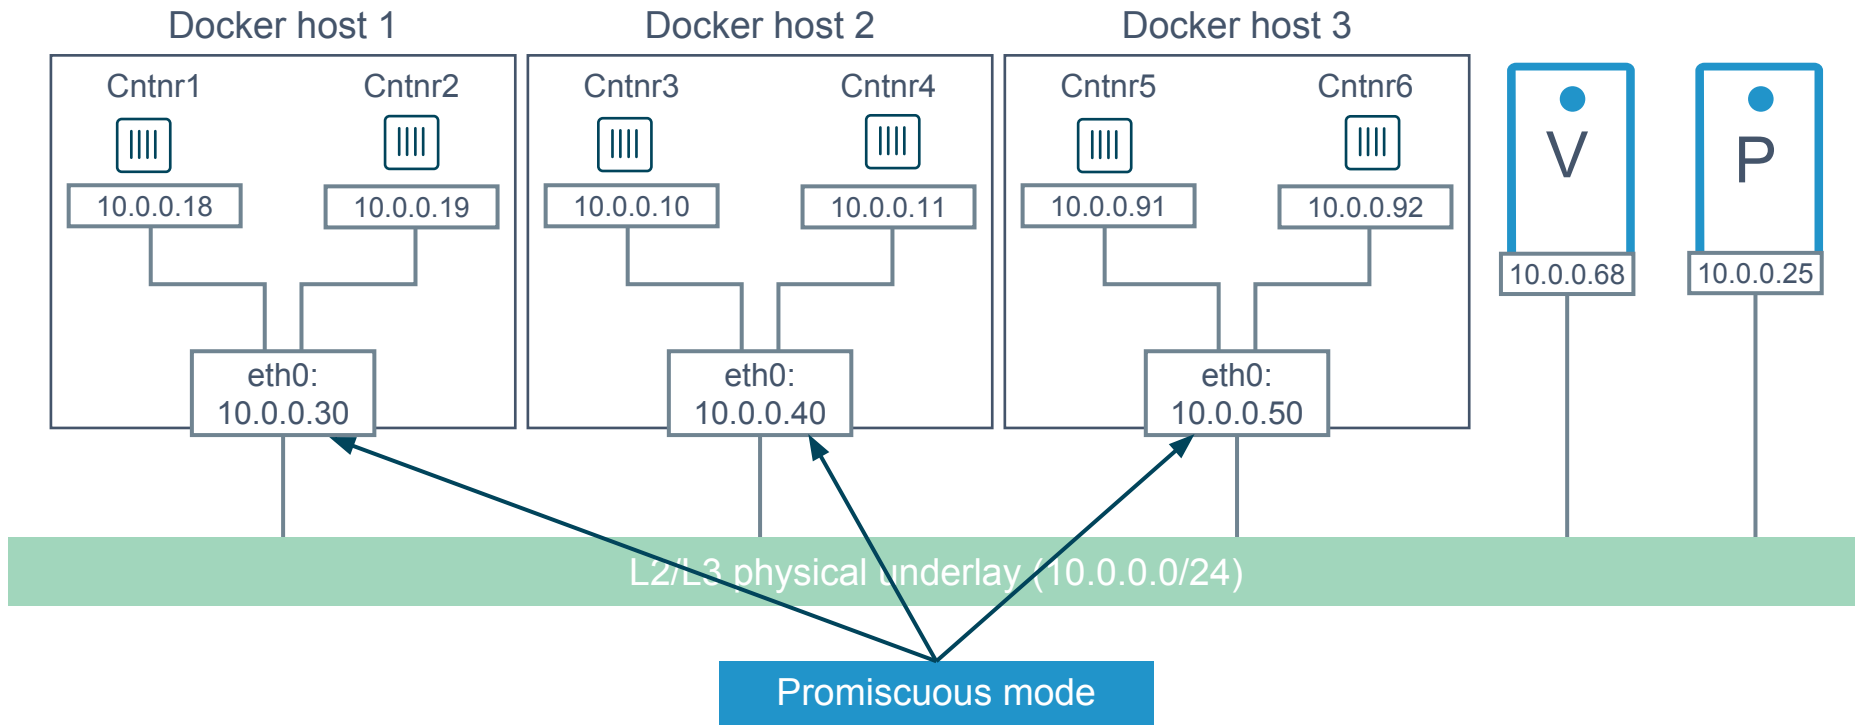
- A way to attach containers to existing networks and VLANs
- Ideal for apps that are not ready to be fully containerized
- Uses the well known MACVLAN Linux network type




What is MACVLAN?



What is MACVLAN?

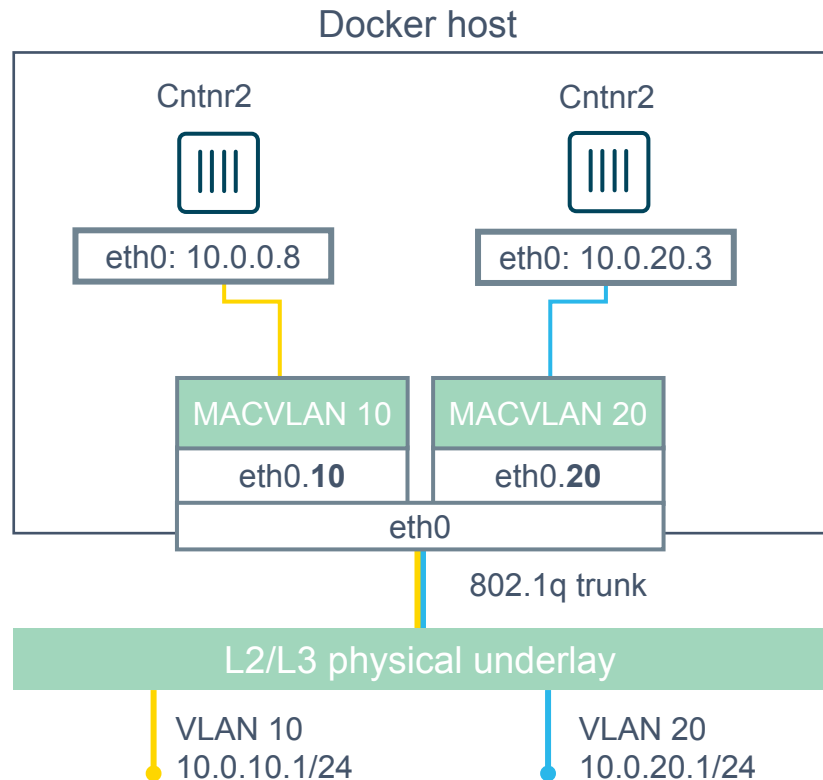


What is MACVLAN?



MACVLAN and Sub-interfaces

- MACVLAN uses **sub-interfaces** to process 802.1Q VLAN tags.
- In this example, two sub-interfaces are used to enable two separate VLANs
- Yellow lines represent VLAN 10
- Blue lines represent VLAN 20



MACVLAN Summary

- Allow containers to be plumbed into existing VLANs
- Ideal for integrating containers with existing networks and apps
- High performance (no NAT or Linux bridge...)
- Every container gets its own **MAC** and **routable IP** on the physical underlay
- Uses **sub-interfaces** for 802.1q VLAN tagging
- Requires **promiscuous mode!**

MACVLAN Summary

- Allow containers to be plumbed into existing VLANs
- Ideal for integrating containers with existing networks and apps
- High performance (no NAT or Linux bridge...)
- Every container gets its own **MAC** and **routable IP** on the physical underlay
- Uses **sub-interfaces** for 802.1q VLAN tagging
- Requires **promiscuous mode!**

Docker Network Design Best Practices

General Networking Design Guidelines

- Make sure to have the required Docker EE TCP/UDP ports open
- Place managers close to each other. Latency can impact raft traffic.
- Pick the right subnet based on the application requirement.
- Dedicate subnets from the underly to be used as overlay subnets
- Monitor your network (rtt, packets drops, tcp retransmits)
- Use Labels!

Networking Reference Architecture

- https://success.docker.com/article/Docker_Reference_Architecture-Designing_Scalable,_Portable_Docker_Container_Networks

Docker Network Troubleshooting

Common Network Issues

Blocked ports, ports required to be open for network mgmt, control, and data plane

Iptables issues

.....

Used extensively by Docker Networking, must not be turned off

List rules with `$ iptables -S`, `$ iptables -S -t nat`

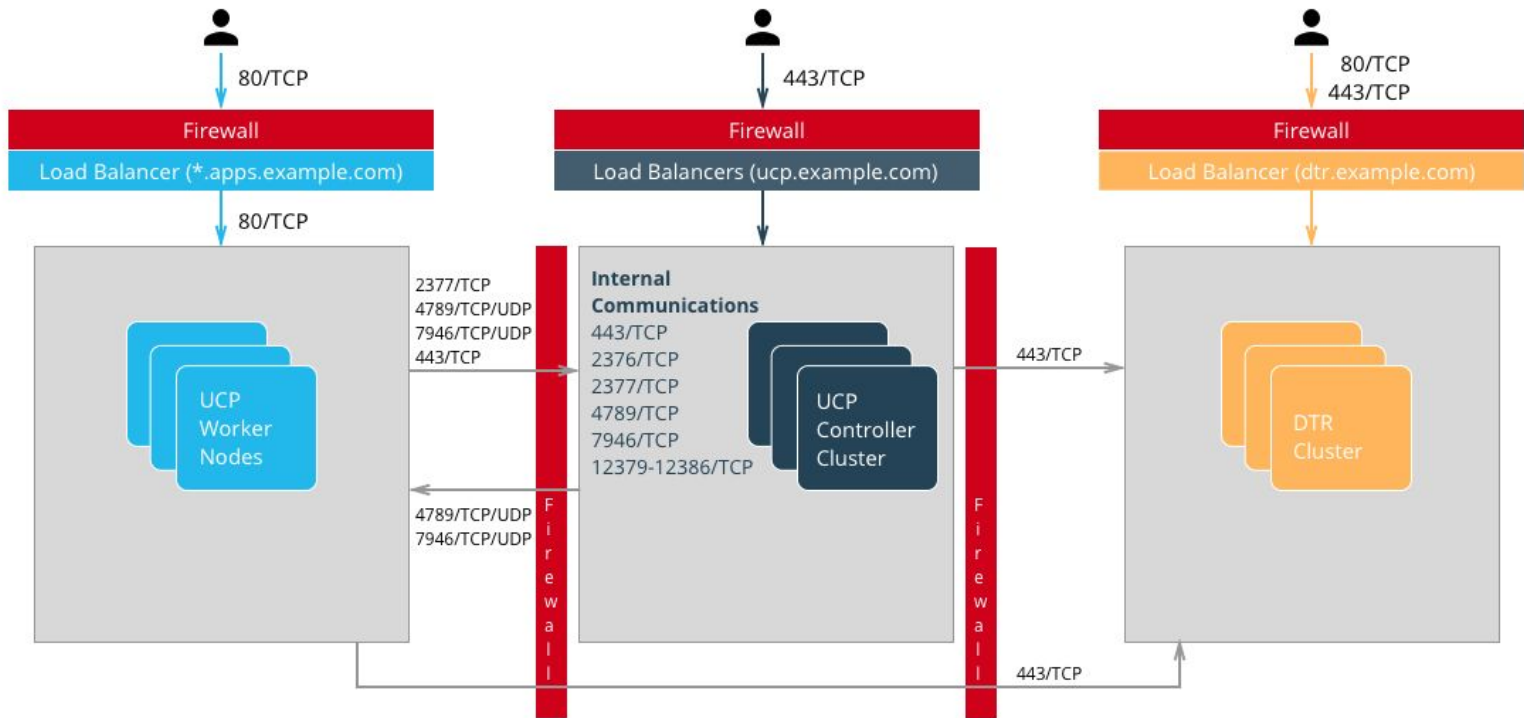
Network state information stale or not being propagated

.....

Destroy and create networks again with same name

General connectivity problems

Required Ports



General Connectivity Issues



Network always gets blamed first :(

Eliminate or prove connectivity first, connectivity can be broken at service discovery or network level



Service Discovery

Test service name resolution or container name resolution

```
drill <service name> (returns  
the service VIP DNS record)  
  
drill tasks.<service name>  
(returns all task DNS records)
```



Network Layer

Test reachability using VIP or container IP

```
task1$ nc -l 5000, task2$  
nc <service ip> 5000  
  
ping <container ip>
```

Netshoot Tool

Has most of the tools you need in a container to troubleshoot common networking problems

```
iperf, tcpdump, netstat, iftop, drill, netcat-openbsd, iproute2,  
util-linux(nsenter), bridge-utils, iputils, curl, ipvsadmin, ethtool...
```

Two Uses

Connect it to a specific **network namespace** (such as a container's) to view the network from that container's perspective

Connect it to a **docker network** to test connectivity on that network

Netshoot Tool

Connect to a container namespace

```
docker run -it --net container:<container_name> nicolaka/netshoot
```

Connect to a network

```
docker run -it --net host nicolaka/netshoot
```

Once inside the **netshoot** container, you can use any of the network troubleshooting tools that come with it

Network Troubleshooting Tools

Capture all traffic to/from port 999 on eth0 on a myservice container

```
docker run -it --net  
container:mymyservice.1.0qlf1kaka0cq38gojf7wcatoa nicolaka/netshoot  
tcpdump -i eth0 port 9999 -c 1 -Xvv
```

See all network connections to a specific task in myservice

```
docker run -it --net  
container:mymyservice.1.0qlf1kaka0cq38gojf7wcatoa nicolaka/netshoot  
netstat -taupn
```

Network Troubleshooting Tools

Test DNS service discovery from one service to another

```
docker run -it --net  
container:myservice.1.bil2mo8inj3r9nyrss1g15qav nicolaka/netshoot drill  
yourservice
```

Show host routing table from inside the netshoot container

```
docker run -it --net host nicolaka/netshoot ip route show
```

Lab 2



THANK YOU